

CMSC 475/675: Introduction to Neural Networks

Review for Exam 1 (Chapters 1, 2, 3, 4)

1. Basics

- Comparison between human brain and von Neumann architecture
- Processing units/nodes (input/output/hidden)
- Activation/node functions (threshold/step, linear-threshold, sigmoid, RBF)
- Network architecture (feed-forward/recurrent nets, layered)
- Connection and weights (excitatory, inhibitory)
- Types of learning (supervised/unsupervised), Hebbian rule,
 - Training samples
 - Overtraining/overfitting problem, cross-validation test

2. Single Layer networks (Perceptron, Adaline, and the delta rule)

- Architecture
- Decision boundary and the problem of linear separability ($w_0 + \sum_{i=1}^n x_i w_i = 0$)
- Perceptron
 - learning rule (only when $x_p \neq d_p : \Delta w_k = \eta \cdot \text{class}(i_p) \cdot i_p$)
 - Perceptron convergence theorem
- **Delta learning rule** and Adaline
 - Error driven: $E = (d_p - \text{net}_p)^2 = (d_p - \sum_k w_k i_{k,p})^2$ or $E = \sum_{p=1}^P (d_p - \text{net}_p)^2 = \sum_{p=1}^P (d_p - \sum_i w_i i_{i,p})^2$
 - Learning rule (delta rule): $\Delta w_k = \eta \cdot (d_p - \text{net}_p) \cdot i_{k,p}$ for each training sample (i_p, d_p)
 - **Gradient descent approach in deriving delta learning rule**
$$\Delta w_k \propto -\partial E / \partial w_k = -2(d_p - \text{net}_p) i_{k,p}$$
 - Local minimum error for gradient descent approach

3. Backpropagation (BP) Networks

- Multi-layer feed-forward architecture with at least one layer of hidden nodes of non-linear and differentiable activation functions
- **Motivation to have non-linear hidden nodes (representational power). Why non-linear?**
- Feed forward computing
- **BP learning**
 - Training samples: (i_p, d_p)
 - Obtain errors at output layer (feed-forward phase): $\delta_k = (d_k - o_k) S'(net_k^{(2)})$
 - **Obtain errors at hidden layer (error backpropagation phase):** $\mu_j = (\sum_k \delta_k w_{k,j}^{(2,1)}) \cdot S'(net_j^{(1)})$
 - Weight update: $\Delta w_{k,j}^{(2,1)} = \eta \cdot \delta_k x_j^{(1)}$, $\Delta w_{j,i}^{(1,0)} = \eta \cdot \mu_j x_i$
 - Why BP learning works (gradient descent to minimize error): $\Delta w_k = -\eta \cdot \partial E / \partial w_k$
 - Learning procedure (batch and sequential modes)
 - In what sense BP learning generalizes the delta rule
- Issues of practical concerns
 - Bias, error bound, training data, initial weights, number and size of hidden layers;
 - Learning rate (momentum, adaptive rate)
- **Advantages and problems with BP learning**
 - Powerful (general function approximator); easy to use; wide applicability; good generalization

- Local minima; overfitting; parameters may be hard to determine; network paralysis; long learning time, black-box; hard to accommodate new samples (non-incremental learning)
- Variations of BP nets
 - Momentum term
 - Adaptive learning rate
 - Quickprop

4. Other Multilayer Nets with Supervised Learning

- Adaptive multilayer nets
 - Why smaller net (with smaller # of hidden nodes) are often preferred
 - Finding “optimal” network size: pruning and growing hidden nodes
- Cascade net (basic ideas):
 - When and how to add a new hidden node
 - What weights are to be trained when a new node is added, and how they are trained
- Prediction networks:
 - BP nets for prediction
 - Recurrent nets: unfolding vs gradient descent
- NN of radial basis function (RBF)
 - Definition of RBF, examples of RBF (especially Gaussian function)
 - Advantages of RBF wrt sigmoid functions
 - RBF for function approximation
- Polynomial networks

Types of questions that may appear on Exam 1:

- **True/False**
 - Backpropagation learning is guaranteed to converge.
- **Definitions**
 - Recurrent networks.
- **Short questions** (conceptual)
 - What are the major differences between human brain and Von Neumann machine?
- **Longer questions**
 - What is the overfitting problem in BP learning? What can you suggest to ease this problem?
- **Apply some NN model to a small concrete problem**
 - Construct a neural network with one hidden node and one output node to solve the XOR problem. The network should be feedforward but not necessarily layered.