

Stack Overflow is a question and answer site for professional and enthusiast programmers. It's 100% free, no registration required.

Take the 2-minute tour ×

## Dealing with Memory Problems in Network with Many Weights

### HANDS-ON CUDA<sup>®</sup> TRAINING IN THE CLOUD

OPENACC · CUDA C/C++ · CUDA FORTRAN · CUDA PYTHON · GPU-ACCELERATED LIBRARIES [LEARN MORE >>](#)

I have a neural network with the architecture 1024, 512, 256, 1 (the input layer has 1024 units, the output layer has 1 unit, etc). I would like to train this network using one of the optimization algorithms in `scipy.optimize`.

The problem is that these algorithms expect the function parameters to be given in one vector; this means that, in my case, I have to unroll all the weights in a vector of length

$$1024*512 + 512*256 + 256*1 = 655616$$

Some algorithms (like `fmin_bfgs`) need to use identity matrices, so they make a call like

```
I = numpy.eye(655616)
```

which, not very surprisingly, produces a `MemoryError`. Is there any way for me to avoid having to unroll all the weights into one vector, short of adapting the algorithms in `scipy.optimize` to my own needs?

[python](#) [numpy](#) [machine-learning](#) [scipy](#)

asked Mar 6 '13 at 10:49



Paul Manta

9,151 4 36 94

Wow, that's a lot of weights to fit. (+1) – [NPE](#) Mar 6 '13 at 10:54

@NPE I'll have to admit that I don't really know what I'm doing. I'm training on images of size 32x32. Should I make the images even smaller? – [Paul Manta](#) Mar 6 '13 at 10:56

Then again, reducing the size of the images to 16x16 and using the architecture 256, 128, 1, I'd still have and unrolled weight vector of length 32896. – [Paul Manta](#) Mar 6 '13 at 11:05

Generally speaking, if you're having performance problems in Python, it might be time to try and do the heavy lifting in a different language. – [jnnnnn](#) Mar 17 '13 at 5:25

[add comment](#)

## 2 Answers

Don't try to fit the weights to a NN using L-BFGS. It doesn't work especially well (see early Yann LeCun papers), and because it's a second-order method you're going to be attempting to approximate the Hessian, which for that many weights is a 655,000 x 650,000 matrix: this introduces a performance overhead that simply won't be justified.

The network isn't that deep: is there a reason you're avoiding standard back-prop? This is just gradient descent if you have access to a library implementation, the gradients are cheap to compute and because it's only a first order method, you'll not have the same performance overhead.

EDIT:

Backpropagation means that the update rule for  $w_i$  at step  $t$  is:

$$w_i(t) = w_i(t-1) - \alpha (dError / dw_i)$$

Also, you've run into the reason why people in vision often use [Convolutional NNs](#): sparse connectivity massively reduces the size of the weight vector where you have one neuron per pixel.

edited Mar 18 '13 at 15:49

answered Mar 18 '13 at 13:01

Ben Allison  
2,521 1 10

+1 For pointing out that the memory problem is caused by the Hessian. I don't avoid standard backprop, I use backprop to calculate the gradient of the error wrt. the network's weights. I still need an optimization algorithm that can traverse the error surface (using the gradient calculated with backprop) to find its minimum. –

[Paul Manta](#) Mar 18 '13 at 14:03

Backprop, at least as I would use the term, means using gradient descent. You can do it stochastically (each example in turn) or in batches. I.e. your update is a step size multiplied by the negative of the gradient. –

[Ben Allison](#) Mar 18 '13 at 14:44

Edited my answer because another useful tidbit occurred to me: convolutional nets are used for this reason. –

[Ben Allison](#) Mar 18 '13 at 15:41

What is the algorithm used to calculate the error gradient wrt. weights called? Also, is there a technical or conceptual reason for "backpropagation" to refer only to training done with gradient descent? – [Paul Manta](#)

Mar 18 '13 at 18:35

I don't think that's an algorithm: the error gradients wrt the weights come from calculus, they're usually analytic (provided your activation functions are differentiable). I think the name backpropagation came from generalising the delta rule, before a lot of research into optimisation methods. It just happens that this is equivalent to gradient descent. Other methods are possible, most notably the Hessian Free method of James Martens, but that's really not trivial to implement and you'll probably get a lot of mileage out of simple backprop. – [Ben Allison](#) Mar 19 '13 at 9:41

show 3 more comments

## HANDS-ON CUDA<sup>®</sup> TRAINING IN THE CLOUD

OPENACC · CUDA C/C++ · CUDA FORTRAN · CUDA PYTHON · GPU-ACCELERATED LIBRARIES [LEARN MORE >>](#)

I think memory-mapping will solve your problem. I suggest using [numpy.memmap](#), which is optimized for numpy arrays. Note that while memory-mapping can be plenty fast, you should keep an eye out for [thrashing](#).

For your general culture, the [mmap](#) module is a more general-purpose memory-mapping library that is included in the standard library.

answered Mar 18 '13 at 9:32

blz  
1,550 13 36

[add comment](#)

Not the answer you're looking for? Browse other questions tagged [python](#) [numpy](#) [machine-learning](#) [scipy](#) or [ask your own question](#).